

OWEN LE

FORWARD DEPLOYED/
APPLIED AI ENGINEER

+447385233956
taivan@hotmail.co.uk
London, United Kingdom
<https://taivanle.github.io/>

PROFILE

Applied AI Engineer deploying production LLM and agent-based systems within complex enterprise environments. Specialised in translating frontier model capabilities into reliable, secure, and measurable solutions that integrate with existing infrastructure. Experienced leading cross-functional builds and owning delivery from architecture through to production deployment.

WORK EXPERIENCE

June 2026 - August 2026 Aptura - Member of Technical Staff

- Designed and implemented Themis, a hybrid agentic grader that benchmarks model outputs against structured rubrics and golden datasets, reducing cost per run by 98.7% (from ~\$25 to \$0.33), cutting evaluation time by ~40%, and improving accuracy by 2 percentage points (from 95% to 97%).
- Built custom evaluation environments (e.g. analyst_code_web and analyst_code_judge) that executed end-to-end grading pipelines, comparing model responses against expected outputs and rubric-defined evaluation criteria.
- Themis combined deterministic validation, rubric-based assessment and LLM-as-a-Judge techniques and agentic grader frameworks to evaluate response quality, correctness and reasoning.
- Developed automated benchmarking workflows to evaluate multiple models consistently across datasets, reducing manual evaluation effort while improving reproducibility.
- Benchmarked Themis against leading evaluation approaches including Gandalf the Grader and LLM-as-a-Judge methodologies, achieving improvements in evaluation accuracy, inference latency and operational cost.
- Worked directly with frontier AI model providers to evaluate production LLMs using large-scale automated grading pipelines.

2023 - June 2026 IBM - Platform & AI Engineer/ FDE

- Delivered AI and automation solutions across 20+ enterprise engagements.
- Led technical builds from proof-of-concept through pilot and production-ready states, defining architecture, integration patterns, and reliability standards.
- Integrated LLM capabilities into enterprise systems via APIs.
- Designed agent-based workflows connecting model reasoning with document repositories, workflow engines, and structured data sources.
- Extended and tested MCP server implementations exposing FileNet (CP4BA on OpenShift) as agent-accessible tools.
- Evaluated frontier LLMs including Claude via API integration, assessing reasoning behaviour and enterprise deployment patterns.
- Led technical discovery sessions, architecture workshops.
- Led AI pilots supporting progression of multi-million-pound enterprise opportunities.

2021 - 2022 Co-Founder & Technical Lead - Blockchain-Based Digital Asset Project

- Co-designed and launched a blockchain-based digital asset collection, contributing to smart contract implementation and on-chain deployment.
- Collaborated on Solidity contract development, token standards, minting logic, wallet integration, and secure on-chain deployment workflows.
- Successfully launched the project and generated revenue through primary digital asset sales.

EDUCATION

2017-2021
UNIVERSITY OF PORTSMOUTH

- Bachelor of Engineering - BEng
- First Class Honours

SKILLS

LLM & Agent Systems: Agent orchestration, tool-calling workflows, enterprise LLM integration, Retrieval Augmented Generation (RAG), Model Context Protocol (MCP), structured prompting, evaluation frameworks, AI governance

AI Frameworks & Vector Databases: LangChain, LangFlow, Milvus, vector search pipelines, embedding workflows

Backend & APIs: Python, TypeScript, JavaScript, SQL, FastAPI, REST APIs, GraphQL, authentication flows

Data & Storage: PostgreSQL, relational data modelling, AI pipeline persistence

Infrastructure & DevOps: Docker, Podman, OpenShift, Kubernetes, Terraform, CI/CD

Frontend: React, HTML, CSS

SELECTED IMPACT

- Designed and built an end to end agentic grader which graded model outputs against rubrics in spreadsheet benchmarking and deck benchmarking.
- Supported enterprise AI engagements across financial services, contributing to multi-million-pound opportunity progression.
- Reduced delivery timelines through reusable AI build frameworks and templated deployment approaches.
- Embedded governance controls into AI workflows, enabling auditable and compliance-aligned deployments.
- Led cross-functional engineering squads delivering production-grade AI pilots within 4-8 week cycles.

LANGUAGES

- English (Native Speaker)